

Delphes for $HH \rightarrow b\bar{b}\tau\tau$ NSBI: full production tuned and validated against CMS

August 26, 2026

Abstract

The full Delphes production is tuned, ntuplised and validated: $3 \times 5\text{M}$ signal events at $\kappa_\lambda = 0, 1, 5$ and 299M $t\bar{t}$ events, 167 GB of analysis ntuples. In the semi-leptonic channel — the one the CMS NSBI result uses — Delphes now reproduces CMS to within a few percent in yield (30 491 vs 28 181 at $\kappa_\lambda = 0$) and agrees in shape on m_{HH} , m_{bb} , p_T^{HH} , ΔR_{bb} , $\Delta\phi_{HH}$, $p_T^{H1,H2}$, $\cos\theta^*$, both τ p_T and p_T^{miss} . Two things were found in the course of this validation and are worth the PI’s attention: the CMS lepton collections were being read *without* ID or isolation, which both broke the semi-leptonic comparison and biased the lepton scale factors (Sec. 2.1); and the b -jet energy *resolution* is untuned by construction, only its scale (Sec. 3.3). One physics residual remains: Delphes makes too many wide-angle τ pairs, visible as an excess at $\Delta R_{\tau\tau} \simeq 2.5\text{--}3.2$ and as high tails on $m_{\tau\tau}$ and m_{vis} (Sec. 3.1). Since the last note the $t\bar{t}$ background has also been compared against its own CMS anchor for the first time and the correction works (Sec. 2.3), and one new defect was found: FastMTT returns no solution for the majority of events, because a Delphes τ_h carries a jet mass far above m_τ (Sec. 3.2).

1 Where we stand

1.1 Production

Everything downstream of Delphes: samples are never re-made, and every correction is a map measured on a private CMS 2024 NanoAODv15 anchor and transferred. The anchor is the PowhegBugFix campaign, the same one the Delphes signal samples were generated from.

sample	events	size	note
signal $\kappa_\lambda = 0, 1, 5$	$3 \times 5.0\text{M}$	7.3 GB	one <code>k1</code> column per event
$t\bar{t}$	298.7M	160.2 GB	1240/1244 shards (Sec. 3.4)
total	313.7M	167.5 GB	33 parquet files

Table 1: Merged ntuples on `/ceph`. Produced by a 1394-job HTCondor campaign, 22.4 min/job, 22.0 GB mean / 26.7 GB peak per job.

The κ_λ value is recorded per event. It is encoded only in the input path and the ntuple schema had no field for it, so a merged signal sample would have been unusable for NSBI — the method is entirely about ratios *between* κ_λ hypotheses. It is now recovered from the production plan at merge time and written as a column.

1.2 Tuning

The τ_h energy correction is the one substantive change since the last note. A multiplicative median scale factor cannot reshape a distribution, and the Delphes τ response carries a one-sided tail that CMS does not have (Fig. 1): at gen $p_T = 20\text{--}30$ GeV one Delphes τ in ten had more than twice the median response, against 1.08 on CMS. The correction now samples the anchor response *distribution* per object, which reproduces the CMS response by construction and fixes scale and width together. m_{vis} , $\tau_1 p_T$, $\tau_2 p_T$ and p_T^{miss} all overlay CMS after this change.

quantity	status	how
b-tag efficiency (<i>b</i> , <i>c</i> , light)	✓ tuned	stochastic re-tag from <code>Jet.Floror</code> at anchor UParT-AK4 Medium
b-jet energy <i>scale</i>	✓ tuned	inverted to unity against <code>GenJet</code> (self-anchored), 0.965–0.989
b-jet energy <i>resolution</i>	✗ tuned	not corrected at all — Sec. 3.3
τ_h ID efficiency	✓ tuned	anchor <code>GenVisTau</code> →DeepTau v2.5 Medium; includes HPS reco inefficiency
τ_h energy response	✓ tuned	<i>distribution</i> resampled from anchor quantiles, not a median ratio
τ_h visible mass	✓ tuned	sampled per jet from anchor per- p_T quantiles
p_T^{miss} resolution	✓ tuned	Gaussian smearing binned in jet H_T
lepton reco efficiency	✓ tuned	per- p_T SFs, now 0.96–1.12 (<i>e</i>), 1.03–1.07 (μ)
di- τ mass estimator	✓ identical	same covariance-free FastMTT on both sides

Table 2: Object status. Every map is derived on the CMS anchor and applied downstream; none needs a card change.




Figure 1: τ_h energy response, Delphes vs the CMS anchor. The Delphes tail above the median is what a multiplicative correction cannot remove and what motivated resampling the distribution instead.

2 Validation against CMS

2.1 The lepton definition — a defect found and fixed

The CMS τ -candidate pool was being built from the raw NanoAOD `Electron` and `Muon` collections, with *no* ID and no isolation requirement. Those collections are loose enough to contain jet fakes that no analysis would use, whereas a Delphes lepton has already passed the card’s efficiency parameterisation, which is effectively an ID. The two pools therefore did not mean the same thing, and the consequences were not cosmetic:

- CMS acquired candidates Delphes cannot produce — 49 950 selected events against Delphes’ 30 491 at $\kappa_\lambda = 0$, a $1.6\times$ excess read at the time as a Delphes acceptance deficit.
- A fake lepton paired with a real τ is back-to-back, so the CMS $\Delta R_{\tau\tau}$ and ΔR_{bb} distributions grew a spurious second peak at $\Delta R \simeq \pi$, plus $m_{\tau\tau}$ and m_{vis} tails and a flattened $\cos\theta^*$.

- The same collections feed `lepton_efficiency`, which sets the `electron_sf` / `muon_sf` anchor targets. The scale factors were therefore tuned toward loose-reco efficiency rather than toward analysis-quality leptons.

The cut is now applied in the NanoAOD *reader*, so the tuning anchor and the validation overlay cannot use different lepton definitions again, and a configured-but-missing ID branch raises rather than silently reverting to the loose collection. Working points are `Electron_mvIso_WP80`, `Muon_tightId` and `Muon_pfRelIso04_all` < 0.15 ; **these should be checked against the CMS $b\bar{b}\tau\tau$ baseline before anything is quoted**. After the fix the lepton fractions match (30% Delphes vs 33% CMS) and the scale factors sit near unity, which is itself the evidence that the two definitions now agree — against the loose collection they had to make up a $\sim 20\%$ gap.

2.2 Result

✓ agree	m_{HH} (the κ_λ observable, at $\kappa_\lambda = 0, 1, 5$); m_{bb} peak; p_T^{HH} ; ΔR_{bb} ; $\Delta\phi_{HH}$; p_T^{H1}, p_T^{H2} ; $\cos\theta^*$; $\tau_1 p_T, \tau_2 p_T$; p_T^{miss}
✗ disagree	$\Delta R_{\tau\tau}$ (Delphes excess at 2.5–3.2), and the $m_{\tau\tau} / m_{\text{vis}}$ high tails that follow from it; m_{bb} <i>width</i>

Table 3: NSBI inputs, semi-leptonic channel, 200k events per κ_λ point against the `PowhegBugFix` anchor including its `ext1` extension. Selected yields Delphes/CMS: 30 491/28 181 ($\kappa_\lambda = 0$), 32 783/31 442 ($\kappa_\lambda = 1$), 24 218/20 443 ($\kappa_\lambda = 5$).

2.3 The $t\bar{t}$ background

The $t\bar{t}$ maps (`maps.ttbar.v1.json`) were applied to every $t\bar{t}$ event in the production but had never been compared back against the anchor they came from. They now have been, on `TTto2L2Nu` with the same selection on both sides, and the correction works: the stock card over-produces τ_h candidates by $1.68\times$ (42 293 against CMS’ 25 241 in the same 200k events) and the re-tag brings that to $1.22\times$. Selected yield moves from 0.70 to 0.89 of CMS.

All ten NSBI features agree in shape both before and after — which is the part to read carefully. The overlays are density-normalised, so a $1.68\times$ error in τ_h *rate* is invisible in them. Shape agreement is not evidence that the τ_h object is right, and the 22% residual after tuning appears in no panel.

3 Open items

3.1 $\Delta R_{\tau\tau}$: too many wide-angle τ pairs

The one physics residual. Delphes has an excess at $\Delta R_{\tau\tau} \simeq 2.5\text{--}3.2$, and the $m_{\tau\tau}$ and m_{vis} high tails are probably the same events counted twice — but that is an inference, not a demonstration, so the argument is worth stating. For legs whose masses are negligible against their energies, $m_{\text{vis}}^2 = 2 p_{T1} p_{T2} (\cosh \Delta\eta - \cos \Delta\phi)$ exactly, and both angular terms grow monotonically as the pair separates; $m_{\tau\tau} = m_{\text{vis}} / \sqrt{x_1 x_2}$ inherits it. A wide pair is therefore a heavy pair — though m_{vis} is monotonic in $\Delta\eta$ and $\Delta\phi$ separately and is *not* a function of ΔR alone, so “wide-angle” rather than “back-to-back” is what ΔR measures. The converse does not follow by itself: a heavy tail could equally come from harder τ , since m_{vis} also grows with $p_{T1} p_{T2}$. What excludes that here is that $\tau_1 p_T$ and $\tau_2 p_T$ both overlay CMS, leaving the angular term as the only place the excess mass can come from. The direct test — cutting $\Delta R_{\tau\tau} < 2.5$ and asking whether the m_{vis} and $m_{\tau\tau}$ tails go with it — has not been run; if the tails survive, these are two problems and not one. This matters for the feature set: dropping $\Delta R_{\tau\tau}$ from the NSBI inputs would remove the most visible

symptom while leaving the identical artefact in $m_{\tau\tau}$ and m_{vis} , which we certainly keep. The cause has to be found rather than worked around.

The leading hypothesis is τ_h mistag: a fake τ from the opposite hemisphere pairs with a real one at $\Delta R \simeq \pi$. This is directly testable — the overlay can split each side into gen-matched and fake τ pairs and report the fake fraction — and has not yet been run at full statistics. The $t\bar{t}$ overlay (Sec. 2.3) now supplies a first measurement of the size of the effect: the stock card makes $1.68\times$ too many τ_h candidates, so there is a real fake excess to account for. It is also actionable if confirmed: the `tau_fake_response` map has its top three p_T bins below the 200-entry threshold and is borrowing them from neighbours, so the high- p_T fake response is extrapolated rather than measured.

3.2 FastMTT returns no solution on the *untuned* samples

Found while preparing the CROWN-matched ntuples, which are built from the *untuned* production. $m_{\tau\tau}$, m_{HH} and p_T^{HH} taken from FastMTT are NaN on 69% of signal and 80% of $t\bar{t}$ events there, while every other branch on those rows is finite. The tuned arm does not have this failure, and the reason is worth stating because it is the clearest example of a map earning its keep.

The mechanism is exact rather than statistical. The hadronic decay weight is non-zero only for $x \geq (m_{\text{vis}}/m_\tau)^2$, with $m_\tau = 1.777$ GeV, so a leg whose visible mass exceeds m_τ has zero likelihood on the whole scan grid and the fit returns nothing. A Delphes τ_h is an AK4 jet and carries the jet mass — 5–15 GeV for a 30–50 GeV jet, three to eight times m_τ , where a real HPS τ_h never exceeds the a_1 mass. Failures therefore track the τ jet mass rather than being a random subset, which is why failing events have a higher visible $m_{\tau\tau}$ (92 vs 72 GeV).

Two consequences. The elliptical signal region evaluates to NaN on those events and rejects them, so it silently doubles as a FastMTT-success filter: a cut CMS quotes at 99% signal efficiency runs at roughly 31% here, on survivors biased in the very mass it selects on. And the failure rate differs between processes (69% vs 80%), so on the untuned samples FastMTT does not cancel between signal and background.

The tuning already fixes exactly this: the `tau_mass` map redraws the τ_h visible mass from the CMS anchor and hard-caps it at 1.70 GeV, just below m_τ , which is the failure that map was written to remove. So this is a limitation of the untuned arm rather than an open defect — but the untuned arm is what the CROWN-matched NSBI inputs are currently built from, so it bites in practice and those inputs should be regenerated from the tuned production or have the cap applied. Two things remain genuinely open: the ellipse should reject a failed fit *explicitly* rather than through $\text{NaN} \rightarrow \infty$, so the two rejection reasons are separable in a cutflow; and capping the mass leaves the jet-level energy scale untouched, which only a τ_h rebuild from the particle-flow constituents would address.

3.3 m_{bb} width: the b -jet resolution is untuned

The b -jet correction inverts the response to unity against Delphes’ *own* GenJet. That fixes the scale, does not touch the width, and never references CMS at all. The τ side, by contrast, now has its full response distribution resampled from the CMS anchor. m_{bb} is therefore the feature receiving the weaker of the two treatments, and its width does not match CMS.

The fix is the direct analogue of what worked for τ : measure the CMS b -jet response as a distribution and resample it. The machinery exists. The cost is the decision — it requires re-deriving the maps and re-running the full campaign (~ 10 h) — and it should be validated on a subset before that is spent. Note also that the card runs without pileup, so part of the CMS b -jet resolution has no counterpart in Delphes to correct.

3.4 Four lost $t\bar{t}$ shards

Four of 1244 $t\bar{t}$ shards (0.3%) could not be produced: the dCache redirector consistently routes those files to a pool node whose certificate does not cover the address it hands out. Bounded retries were added and confirm the failure is deterministic, not transient, so this is a site issue to report rather than a pipeline defect. Shards group input files by size with no relation to physics, so the loss costs MC statistics and introduces no bias; the shortfall is recorded in the merged manifest rather than hidden.

4 Decisions and next steps

1. **Confirm the lepton working points** against the CMS $b\bar{b}\tau\tau$ baseline. Everything downstream follows from them, and the current values are conventional choices rather than the analysis's.
2. **Clamp the τ_h mass in FastMTT** (Sec. 3.2). One line, and it unblocks every $m_{\tau\tau}$ -based quantity; until it is in, use the visible mass.
3. **Run the gen-matched/fake split at full statistics** to test whether the $\Delta R_{\tau\tau}$ excess is τ mistag. Cheap, and it decides whether Sec. 3.1 is a fixable map problem or something deeper.
4. **Decide whether to tune the b -jet resolution** (Sec. 3.3). This is the only item costing a re-production, so it should be validated on a subset first.
5. **Then train.** The ntuples are ready and κ_λ -labelled; the feature set should be settled by measured κ_λ separation and by training with and without candidate features, not by inspecting one-dimensional marginals.

5 Known limitations

Accepted by design. FastMTT is covariance-free and is run identically on both sides. The cancellation argument holds for the tuned samples, where `tau_mass` keeps the τ_h leg below m_τ . It does *not* hold for the untuned ones: Sec. 3.2 shows the fit there fails at different rates for signal and $t\bar{t}$, so the failures do not divide out and the visible mass should be used instead. The card runs without pileup, so jet energy resolution is untuned and pileup jets cannot exist; neither is fixable downstream. p_T^{miss} smearing is now binned in jet H_T but remains isotropic where the real resolution is anisotropic.

A scope bound, not a defect. `tau_mistag` is measured on the b -rich signal anchor and fake rates are strongly jet-flavour dependent, so it is not valid for fake-dominated backgrounds (QCD, W+jets). The present signal study is unaffected; a background-heavy analysis would need those measured separately. There is likewise no CMS anchor for DY, so DY cannot currently be tuned.

Untested — risks, not accepted limitations. The sampled τ_h mass is drawn independently of the τ energy and of the partner leg, whereas the true visible mass correlates with decay mode; size unknown. Every map is flat-extrapolated above 300 GeV. That is probably benign here — only $\sim 2.8\%$ of b -jets exceed it, the card's b -tag formula tends to a constant so flat is the right asymptote, and κ_λ discrimination lives near threshold ($m_{HH} \simeq 250\text{--}400$ GeV, where the triangle's s -channel propagator makes the interference most κ_λ -sensitive) rather than in the box-dominated tail. A boosted or high- m_{HH} category would change that, by making the extrapolated region the signal region.